

```
In [0]: import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras import backend as K
import pickle
```

```
In [0]: # Check that TF 2.1.0 is in use
# I use colab and accelerate by their GPU so tf is 2.2.0
print(tf.__version__)
```

2.2.0


**Points awarded for correct working models, questions, and plots.**

[+10 per model correct and working -5 for failure on either]

## 1.1 Answer the following questions:

### 1. Explain the indication of overfitting and how this occurs (provide plot supporting your answer)? [+4 answer, +4 plot, +2 answer and plot agree]

The **indication** of overfitting is that when our model does much better on the training set than on the test/validation set, then we're likely overfitting. Just like the figure below, we can see that the accuracy of the training dataset is higher than of the validation set and the loss on the validation set goes much higher then on the training set, that means the performance of our model is only good on the training set and it can be overfitting.



The **reason** for the overfitting is that: when the algorithm is too complicated or flexible but the training data is limited, it can end up “memorizing the training dataset” instead of finding the features, so that the model will make predictions based on the relationship between inputs and labels in the training dataset and perform unusually well on its training data but badly on the validation set, then we can see a large difference of the curves in this figure.

### 2. Explain how overfit can hinder performance of a model when deployed. [+6 answer]

Just as explained in problem1.1.1, overfitting happens when the algorithm is too complicated but the training data is limited and leads to an unusually good performance on the training set but may fail to fit additional data, and this may affect the accuracy of predicting future observations. So when we deployed this kind of overfitting model to the real test case, the output accuracy maybe not very good and the performance may not be very robust.

### 3. Name two ways to avoid this. [+2 answer, +2 answer]

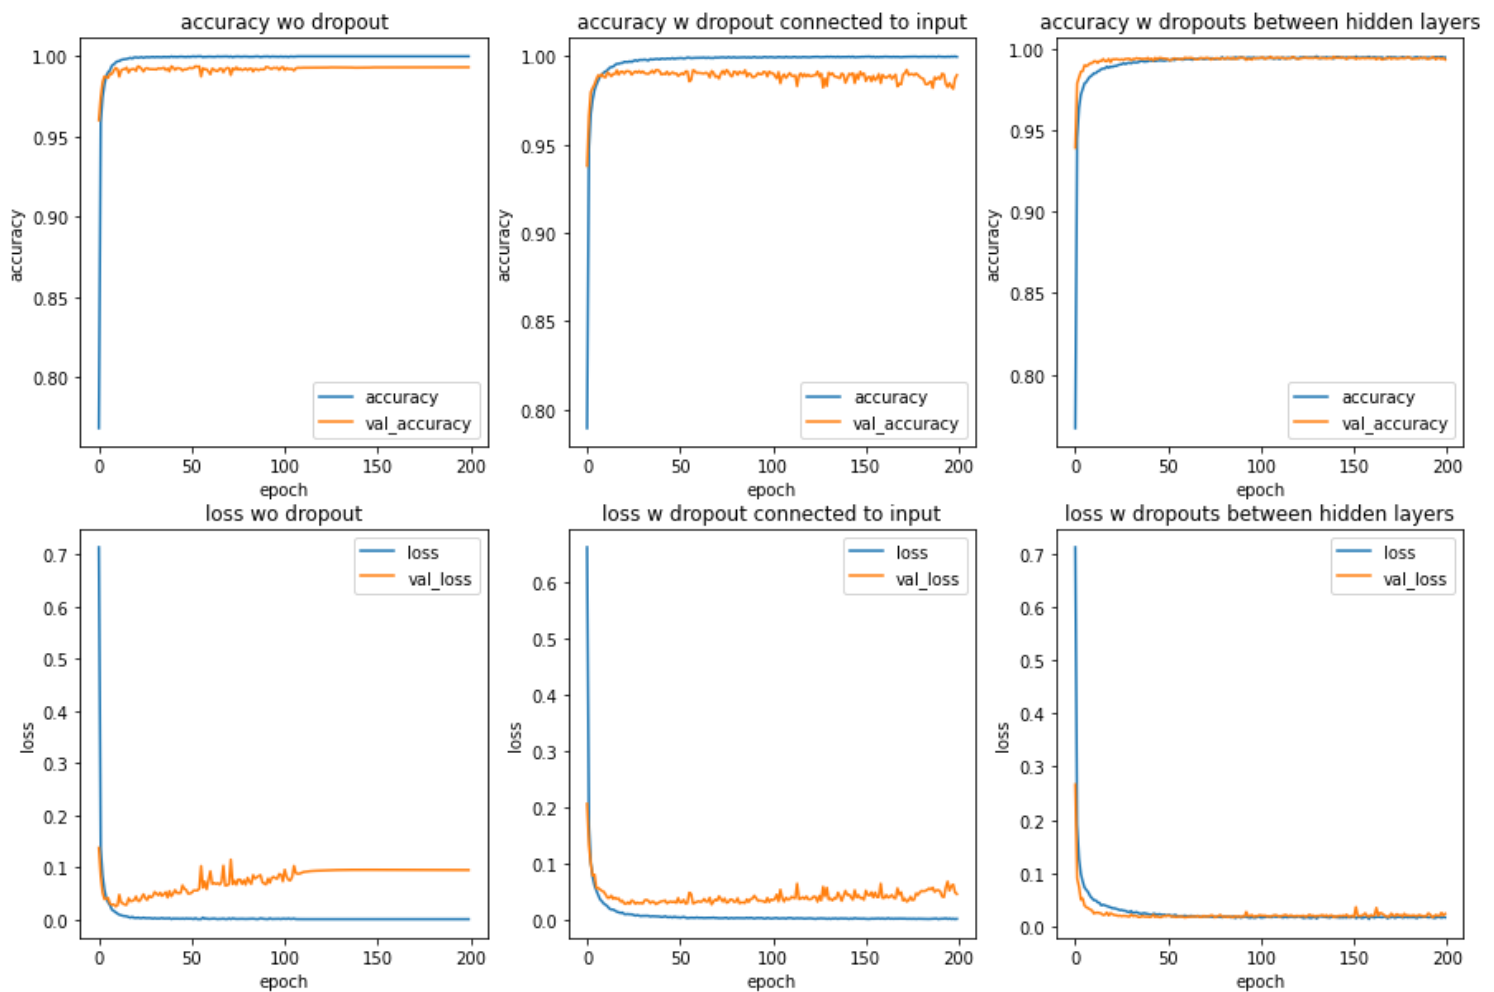
1. We can train our model with more data to prevent overfitting and add more diversity to the training set.
2. We can do the data simplification. Since the overfitting can occur due to the complexity of a model, we can try to decrease the complexity of the model to make it simple enough that it does not overfit. We can achieve that by pruning a decision tree, reducing the number of parameters in a neural network, and using dropout on a neutral network.

## 1.2 Answer the following question:

### 1. Explain how dropout affected your loss (provide plot supporting your answer). [+5 answer, +5 plot]

Here I did two experiments according to the requirement.

1. The first one is that a dropout of 30% was inserted between the input and the first hidden layer. Then the model was trained again and the final accuracy result is **98.92%**. Compared with the accuracy score without the dropout layer(**98.32%**), the accuracy decreases a little. **The second column in the figure below shows the accuracy and loss change during the training process under this case, and we can see that the overfitting situation did not improve too much and the accuracy curve even became worse.**
2. The second experiment is that the dropout between input and hidden was removed and a dropout to each hidden layer except between softmax and output layer was added. The final accuracy score is **98.37%**, better than the result in the first experiment and the result without dropout. **From the third column of the figure below, we can see that the curves of the training set and validation set are quite similar no matter from the accuracy plot or from the loss plot, which means that the overfitting problem was solved.**



The **reason** for this phenomenon is that the dropout actually is trying to ignore units (i.e., neurons) during the training phase of a certain set of neurons which is chosen at random, the ignoring units are not considered during a particular forward or backward pass, so it is a kind of data simplification and can help to overcome the overfitting. However, if the dropout layer is inserted between the input layer and the first hidden layer, the input space information can be broken and some information from this raw input can be ignored, which may reduce the performance of the model and results to a lower accuracy(just as the result shown in the second column). So we need to avoid doing that and only apply the dropout between other layers, then we can see a good result and solve the overfitting problem.



**Bonus Answer the following question:**

**1. Considering that encoder and decoder can be constructed as separate components, trained as a single unit, and then separated for use . What uses can you brainstorm? [+5 bonus makeup points]**



From my perspective, I think the separated autoencoder can be used in the encryption field like **secure communication, image encryption** ...

Say for example, when we need to do the secure communication, we can train an autoencoder first and then separated the model, and the sender can keep the encoder part and input the message into the encoder then send the output of the encoder. When the receiver gets the signal, the raw message can be recovered by the decoder. I think compared with the other secure communication method, use encoder and decoder as the key can be more secure, but the accuracy may be the main limitation. Probably an overfitting model can be better for this application. Also, the image encryption is similar, we can train AE first use the encoder and decoder as the keys.

2.1 Linear AE points for constructed model, no questions here.



2.2 Convolutional AE points for constructed model, no questions here.




**2.3 Report histogram plot, mean and std. dev. of normal data, and confusion matrix for 2 standard deviations as results. Discuss your loss plot.** [+10 for greater than 75 TP, +10 all else]

Since the model performance in problem2.2 is good, I design a model based on the Convolutional AE in problem2.2 and further increase the number of convolutional kernels to faced with the larger input. The details of this model are shown in the below figure. For the encoder part, this model has three conv2d layers and three max-pooling layers, for the decoder part, this model has three conv2d layers and three up\_sampling layers. For the number of convolutional kernels in each conv2d layer, I used to try **16-8-8-16**, **32-16-16-32**, **32-4-4-32** and **32-16-4-4-16-32**. Finally, I choose **32-16-4-4-16-32** since this one has the best performance and when the models are initialized randomly, this one can always get a stable result. Also, the overfitting problem is not very obvious for this combination, so I think the 32-16-4-4-16-32 has a good balance of complexity and performance.

Model: "sequential\_23"

Layer (type)	Output Shape	Param #
conv2d_131 (Conv2D)	(None, 64, 312, 32)	320
max_pooling2d_55 (MaxPooling)	(None, 32, 156, 32)	0
dropout_6 (Dropout)	(None, 32, 156, 32)	0
conv2d_132 (Conv2D)	(None, 32, 156, 16)	4624
max_pooling2d_56 (MaxPooling)	(None, 16, 78, 16)	0
conv2d_133 (Conv2D)	(None, 16, 78, 4)	580
max_pooling2d_57 (MaxPooling)	(None, 8, 39, 4)	0
conv2d_134 (Conv2D)	(None, 8, 39, 4)	148
up_sampling2d_53 (UpSampling)	(None, 16, 78, 4)	0
conv2d_135 (Conv2D)	(None, 16, 78, 16)	592
up_sampling2d_54 (UpSampling)	(None, 32, 156, 16)	0
conv2d_136 (Conv2D)	(None, 32, 156, 32)	4640
up_sampling2d_55 (UpSampling)	(None, 64, 312, 32)	0
conv2d_137 (Conv2D)	(None, 64, 312, 1)	289
Total params: 11,193		
Trainable params: 11,193		
Non-trainable params: 0		
None		

The overfitting problem for this model is not very serious, so I just add one dropout layer and then we can see two very similar loss curves from the figure below. **I think the reason for the less overfitting is that convolutional layers just have few parameters, they need less regularization to begin with. Furthermore, the max-pooling layers can also take the roles of reducing the complexity of the model.**

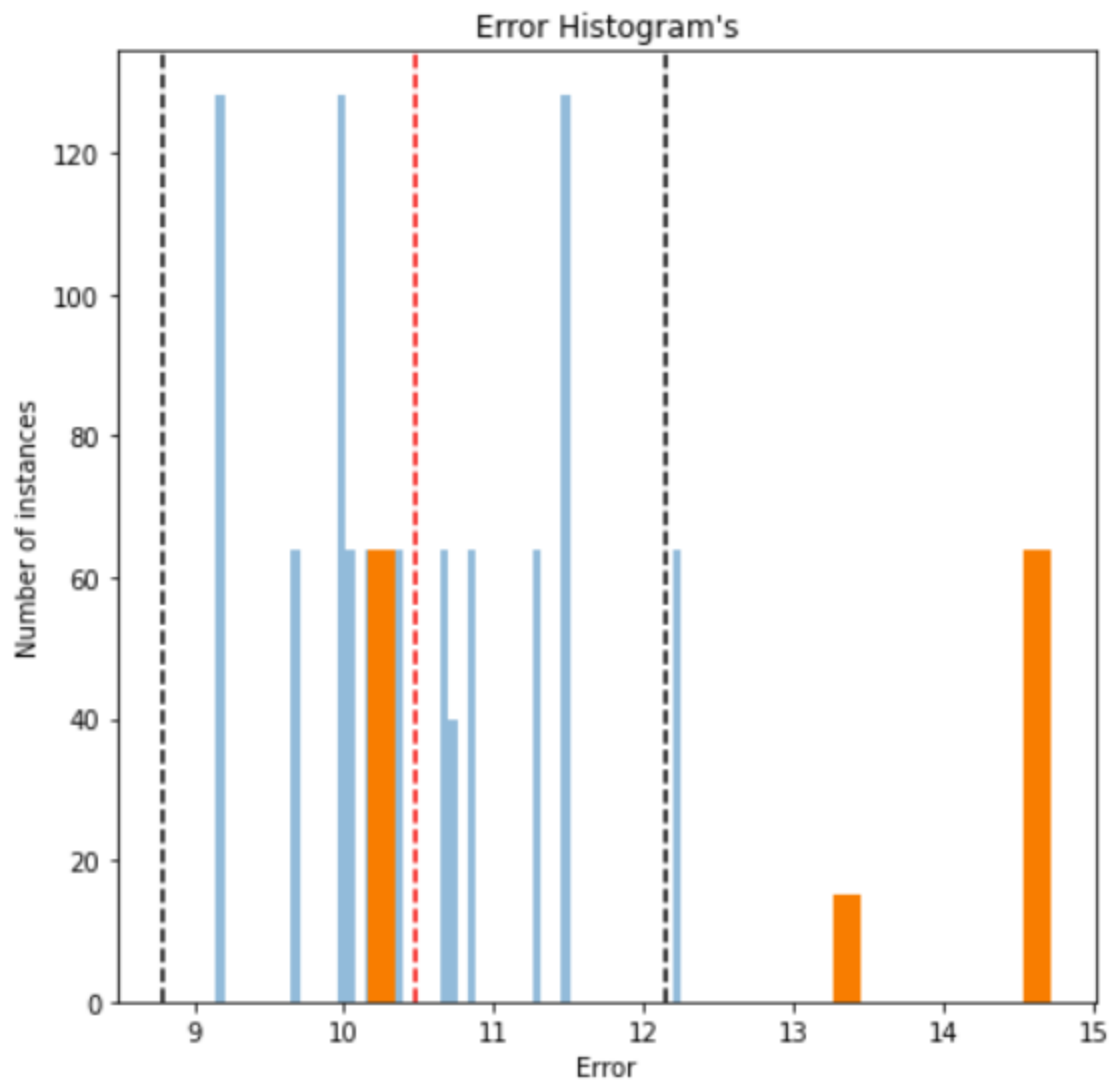


To justify my model beyond the result: I ran it **repeatedly with random initializations for five times** and still receive the **same result(TP=79)**. So I do not think I get a lucky initialization in the weight space and the results are very stable. For the non-linearity I also try the sigmoid but I didn't get any good results based on that. I think the sigmoid function may be more useful in classification problems since the output is scaled nonlinearly but it may not be a good choice for image reconstruction. I trained my model for 150 and 80 epochs and I think **80 epochs** is better due to the similar performance and better time consumption. For anything in the data that affected model building, I think the most important thing is the **shape of the input data**. It can not only affect the input layers' shape but also affect how many max-pooling layers we can use. I eventually just used **one channel** from the input data because I noticed that one channel is more effective than others. I also try to combine eight channels data together and make an input data set like (1000,64,3138,1), but the results did not make any sense, so I just use one channel. **I think that means some time wash the data before the training process can significantly benefit the results.**



The mean of normal data is 10.4711

and standard deviation is 0.8402



**Reminder: Achieve better than 75 anomalies**

```
In [146]: #Import dataset and normalize to [0,1]
mnist = tf.keras.datasets.mnist
(data_train, labels_train), (data_test, labels_test) = mnist.load_data()
#Normalize
data_train = data_train / 255
data_test = data_test / 255
#Reshape
data_train = data_train.reshape(data_train.shape[0], 28, 28, 1)
data_test = data_test.reshape(data_test.shape[0], 28, 28, 1)

#Create labels as one-hot vectors
labels_train = tf.keras.utils.to_categorical(labels_train, num_classes=10)
labels_test = tf.keras.utils.to_categorical(labels_test, num_classes=10)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
11493376/11490434 [=====] - 0s 0us/step

## Section 1 - CNN's

Fill in the model:

- Input: 28x28x1 grayscale image (1 specifies single channel grayscale).
- 1st hidden: 2D convolutional layer with 256 feature maps and 3x3 filters.
- 2nd hidden: A 2x2 maxpool layer.
- 3rd hidden: 2D convolutional layer with 128 feature maps and 3x3 filters.
- 4th hidden: A 2x2 maxpool layer.
- 5th hidden: Flatten layer to map 2D to 1D vector.
- 6th hidden: Dense layer of 100 perceptrons.
- 7th hidden: Dense layer of 100 perceptrons.
- Output: 10 perceptrons for classification.

**Activations, bias, loss function, and optimizer are your choice.**

**Train for 200 epochs**

### 1.1 Overfitting

```

In [147]: #Create and train model architecture
def CNN_overfit():

    #Easiest way to build model in Keras is using Sequential. It allows model to be build layer by layer as we will do here
    model = tf.keras.models.Sequential()

    #### Fill in Model ####
    model.add(tf.keras.layers.Conv2D(256, kernel_size=(3, 3), activation='relu', input_shape=(28, 28, 1)))
    model.add(tf.keras.layers.MaxPooling2D(pool_size=(2, 2)))
    model.add(tf.keras.layers.Conv2D(128, (3, 3), activation='relu'))
    model.add(tf.keras.layers.MaxPooling2D(pool_size=(2, 2)))

    model.add(tf.keras.layers.Flatten())
    model.add(tf.keras.layers.Dense(100, activation='relu'))
    model.add(tf.keras.layers.Dense(100, activation='relu'))
    model.add(tf.keras.layers.Dense(10, activation='softmax'))

    return model

CNN_overfit = CNN_overfit()
print(CNN_overfit.summary())
CNN_overfit.compile(loss='categorical_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
history_overfit = CNN_overfit.fit(data_train, labels_train, validation_data=(data_test, labels_test), epochs=10, batch_size=1000, shuffle=True)
scores = CNN_overfit.evaluate(data_test, labels_test)
print("Accuracy: %.2f%%" %(scores[1]*100))

```

Model: "sequential\_25"

Layer (type)	Output Shape	Param #
conv2d_140 (Conv2D)	(None, 26, 26, 256)	2560
max_pooling2d_60 (MaxPooling)	(None, 13, 13, 256)	0
conv2d_141 (Conv2D)	(None, 11, 11, 128)	295040
max_pooling2d_61 (MaxPooling)	(None, 5, 5, 128)	0
flatten_1 (Flatten)	(None, 3200)	0
dense_3 (Dense)	(None, 100)	320100
dense_4 (Dense)	(None, 100)	10100
dense_5 (Dense)	(None, 10)	1010

Total params: 628,810

Trainable params: 628,810

Non-trainable params: 0

None

Epoch 1/10

60/60 [=====] - 12s 205ms/step - loss: 0.6882 - accuracy: 0.7825 - val\_loss: 0.3163 - val\_accuracy: 0.8999

Epoch 2/10

60/60 [=====] - 12s 204ms/step - loss: 0.1432 - accuracy: 0.9552 - val\_loss: 0.0772 - val\_accuracy: 0.9752

Epoch 3/10

60/60 [=====] - 12s 204ms/step - loss: 0.0811 - accuracy: 0.9746 - val\_loss: 0.0732 - val\_accuracy: 0.9759

Epoch 4/10

60/60 [=====] - 12s 204ms/step - loss: 0.0562 - accuracy: 0.9829 - val\_loss: 0.0359 - val\_accuracy: 0.9887

Epoch 5/10

60/60 [=====] - 12s 204ms/step - loss: 0.0402 - accuracy: 0.9876 - val\_loss: 0.0318 - val\_accuracy: 0.9894

Epoch 6/10

60/60 [=====] - 12s 204ms/step - loss: 0.0316 - accuracy: 0.9899 - val\_loss: 0.0478 - val\_accuracy: 0.9838

Epoch 7/10

60/60 [=====] - 12s 204ms/step - loss: 0.0259 - accuracy: 0.9919 - val\_loss: 0.0273 - val\_accuracy: 0.9918

Epoch 8/10

60/60 [=====] - 12s 204ms/step - loss: 0.0199 - accuracy: 0.9937 - val\_loss: 0.0238 - val\_accuracy: 0.9924

Epoch 9/10

60/60 [=====] - 12s 204ms/step - loss: 0.0162 - accuracy: 0.9950 - val\_loss: 0.0269 - val\_accuracy: 0.9913

Epoch 10/10

60/60 [=====] - 12s 204ms/step - loss: 0.0126 - accuracy: 0.9960 - val\_loss: 0.0276 - val\_accuracy: 0.9912



313/313 [=====] - 2s 5ms/step - loss: 0.0276 - accuracy: 0.9912

Accuracy: 99.12%

```
In [0]: # Information contained in history dict.  
print(history_overfit.history.keys())
```

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
In [0]: #acc scores = 99.32%  
#Plot accuracy vs epoch  
plt.figure(figsize=(15,7))  
plt.subplot(121)  
plt.plot(history_overfit.history['accuracy'], label='accuracy')  
plt.plot(history_overfit.history['val_accuracy'], label='val_accuracy')  
plt.legend(loc='lower right')  
plt.title('accuracy')  
plt.xlabel('epoch')  
plt.ylabel('acc')  
#### Fill in plot ####  
  
#Plot loss vs epoch  
plt.subplot(122)  
plt.plot(history_overfit.history['loss'], label='loss')  
plt.plot(history_overfit.history['val_loss'], label='val_loss')  
plt.legend(loc='upper right')  
plt.title('loss')  
plt.xlabel('epoch')  
plt.ylabel('loss')  
plt.show()  
#### Fill in plot ####
```



## 1.2 Improvements

Using the network above, (1) insert a dropout of 30% between the input and first hidden layer. Run the model again and make note of the result. Next, (2) remove the dropout between input and hidden and add a dropout to each hidden layer except between softmax and output layer. Plot accuracy and loss only for (2). What do you observe for (2)?.

For 1.2 (1):

```
In [148]: def CNN_dropout_hidden():
    model = tf.keras.models.Sequential()
    #### Fill in model ####
    model.add(tf.keras.layers.Input(shape=(28, 28, 1)))
    model.add(tf.keras.layers.Dropout(0.3))
    model.add(tf.keras.layers.Conv2D(256, kernel_size=(3, 3), activation='relu'))
    model.add(tf.keras.layers.MaxPooling2D(pool_size=(2, 2)))
    model.add(tf.keras.layers.Conv2D(128, (3, 3), activation='relu'))
    model.add(tf.keras.layers.MaxPooling2D(pool_size=(2, 2)))
    model.add(tf.keras.layers.Flatten())
    model.add(tf.keras.layers.Dense(100, activation='relu'))
    model.add(tf.keras.layers.Dense(100, activation='relu'))
    model.add(tf.keras.layers.Dense(10, activation='softmax'))

    return model

#Compile and train the model
CNN_dropout_hidden = CNN_dropout_hidden()
CNN_dropout_hidden.compile(loss='categorical_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
history_dropout_hidden = CNN_dropout_hidden.fit(data_train, labels_train, validation_data=(data_test, labels_test), epochs=10, batch_size=1000, shuffle=True)
scores_dropout_hidden = CNN_dropout_hidden.evaluate(data_test, labels_test)
print("Accuracy: %.2f%%" %(scores_dropout_hidden[1]*100))
```

Epoch 1/10

60/60 [=====] - 12s 205ms/step - loss: 0.7257 - accuracy: 0.7597 - val\_loss: 0.2435 - val\_accuracy: 0.9432

Epoch 2/10

60/60 [=====] - 12s 204ms/step - loss: 0.1565 - accuracy: 0.9508 - val\_loss: 0.1395 - val\_accuracy: 0.9606

Epoch 3/10

60/60 [=====] - 12s 204ms/step - loss: 0.1022 - accuracy: 0.9673 - val\_loss: 0.0919 - val\_accuracy: 0.9771

Epoch 4/10

60/60 [=====] - 12s 204ms/step - loss: 0.0748 - accuracy: 0.9766 - val\_loss: 0.0908 - val\_accuracy: 0.9830

Epoch 5/10

60/60 [=====] - 12s 204ms/step - loss: 0.0584 - accuracy: 0.9821 - val\_loss: 0.0690 - val\_accuracy: 0.9840

Epoch 6/10

60/60 [=====] - 12s 204ms/step - loss: 0.0488 - accuracy: 0.9850 - val\_loss: 0.0657 - val\_accuracy: 0.9875

Epoch 7/10

60/60 [=====] - 12s 204ms/step - loss: 0.0411 - accuracy: 0.9865 - val\_loss: 0.0827 - val\_accuracy: 0.9871

Epoch 8/10

60/60 [=====] - 12s 204ms/step - loss: 0.0350 - accuracy: 0.9889 - val\_loss: 0.0577 - val\_accuracy: 0.9867

Epoch 9/10

60/60 [=====] - 12s 204ms/step - loss: 0.0300 - accuracy: 0.9899 - val\_loss: 0.0394 - val\_accuracy: 0.9905

Epoch 10/10



60/60 [=====] - 12s 204ms/step - loss: 0.0279 - accuracy: 0.9907 - val\_loss: 0.0563 - val\_accuracy: 0.9878

313/313 [=====] - 2s 5ms/step - loss: 0.0563 - accuracy: 0.9878

Accuracy: 98.78%

```
In [0]: # acc scores = 98.92%
# Plot accuracy vs epoch
plt.figure(figsize=(15,7))
plt.subplot(121)
plt.plot(history_dropout_hidden.history['accuracy'], label='accuracy')
plt.plot(history_dropout_hidden.history['val_accuracy'], label='val_accuracy')
plt.legend(loc='lower right')
plt.title('accuracy')
plt.xlabel('epoch')
plt.ylabel('acc')
#### Fill in plot ####

#Plot loss vs epoch
plt.subplot(122)
plt.plot(history_dropout_hidden.history['loss'], label='loss')
plt.plot(history_dropout_hidden.history['val_loss'], label='val_loss')
plt.legend(loc='upper right')
plt.title('loss')
plt.xlabel('epoch')
plt.ylabel('loss')
#### Fill in plot ####
plt.show()
```



For 1.2 (2):

```

In [149]: #Create and train model architecture
def CNN_dropout_hidden2():
    model = tf.keras.models.Sequential()
    #### Fill in model ####
    model.add(tf.keras.layers.Conv2D(256, kernel_size=(3, 3), activation='relu', input_shape=(28, 28, 1)))
    model.add(tf.keras.layers.Dropout(0.3))
    model.add(tf.keras.layers.MaxPooling2D(pool_size=(2, 2)))
    model.add(tf.keras.layers.Dropout(0.3))
    model.add(tf.keras.layers.Conv2D(128, (3, 3), activation='relu'))
    model.add(tf.keras.layers.Dropout(0.3))
    model.add(tf.keras.layers.MaxPooling2D(pool_size=(2, 2)))
    model.add(tf.keras.layers.Dropout(0.3))
    model.add(tf.keras.layers.Flatten())
    model.add(tf.keras.layers.Dropout(0.3))
    model.add(tf.keras.layers.Dense(100, activation='relu'))
    model.add(tf.keras.layers.Dropout(0.3))
    model.add(tf.keras.layers.Dense(100, activation='relu'))
    model.add(tf.keras.layers.Dropout(0.3))
    model.add(tf.keras.layers.Dense(10, activation='softmax'))

    return model

#Compile and train the model
CNN_dropout_hidden2 = CNN_dropout_hidden2()
CNN_dropout_hidden2.compile(loss='categorical_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
history_dropout_hidden2 = CNN_dropout_hidden2.fit(data_train, labels_train, validation_data=(data_test, labels_test), epochs=10, batch_size=1000, shuffle=True)
scores_dropout_hidden2 = CNN_dropout_hidden2.evaluate(data_test, labels_test)
print("Accuracy: %.2f%%" %(scores_dropout_hidden2[1]*100))

```




Epoch 1/10  
60/60 [=====] - 15s 258ms/step - loss: 0.7937 - accuracy: 0.7398 - val\_  
loss: 0.2112 - val\_accuracy: 0.9544  
Epoch 2/10  
60/60 [=====] - 15s 256ms/step - loss: 0.2062 - accuracy: 0.9398 - val\_  
loss: 0.1037 - val\_accuracy: 0.9762  
Epoch 3/10  
60/60 [=====] - 15s 256ms/step - loss: 0.1370 - accuracy: 0.9596 - val\_  
loss: 0.0825 - val\_accuracy: 0.9800  
Epoch 4/10  
60/60 [=====] - 15s 256ms/step - loss: 0.1051 - accuracy: 0.9702 - val\_  
loss: 0.0549 - val\_accuracy: 0.9872  
Epoch 5/10  
60/60 [=====] - 15s 256ms/step - loss: 0.0888 - accuracy: 0.9744 - val\_  
loss: 0.0538 - val\_accuracy: 0.9881  
Epoch 6/10  
60/60 [=====] - 15s 256ms/step - loss: 0.0766 - accuracy: 0.9777 - val\_  
loss: 0.0442 - val\_accuracy: 0.9883  
Epoch 7/10  
60/60 [=====] - 15s 256ms/step - loss: 0.0663 - accuracy: 0.9804 - val\_  
loss: 0.0376 - val\_accuracy: 0.9909  
Epoch 8/10  
60/60 [=====] - 15s 256ms/step - loss: 0.0611 - accuracy: 0.9822 - val\_  
loss: 0.0338 - val\_accuracy: 0.9915  
Epoch 9/10  
60/60 [=====] - 15s 256ms/step - loss: 0.0578 - accuracy: 0.9834 - val\_  
loss: 0.0310 - val\_accuracy: 0.9926  
Epoch 10/10  
60/60 [=====] - 15s 256ms/step - loss: 0.0561 - accuracy: 0.9839 - val\_  
loss: 0.0300 - val\_accuracy: 0.9924  
313/313 [=====] - 2s 5ms/step - loss: 0.0300 - accuracy: 0.9924  
Accuracy: 99.24%

```

In [0]: # acc scores = 99.37%
# Plot accuracy vs epoch
plt.figure(figsize=(15,7))
plt.subplot(121)
plt.plot(history_dropout_hidden2.history['accuracy'], label='accuracy')
plt.plot(history_dropout_hidden2.history['val_accuracy'], label='val_accuracy')
plt.legend(loc='lower right')
plt.title('accuracy')
plt.xlabel('epoch')
plt.ylabel('acc')
#### Fill in plot ####

#Plot loss vs epoch
plt.subplot(122)
plt.plot(history_dropout_hidden2.history['loss'], label='loss')
plt.plot(history_dropout_hidden2.history['val_loss'], label='val_loss')
plt.legend(loc='upper right')
plt.title('loss')
plt.xlabel('epoch')
plt.ylabel('loss')
#### Fill in plot ####
plt.show()

```



Summary for 1.2 Dropout improvements

```
In [0]: plt.figure(figsize=(15,10))
```

```
# Plot accuracy vs epoch
```

```
plt.subplot(231)
plt.plot(history_overfit.history['accuracy'], label='accuracy')
plt.plot(history_overfit.history['val_accuracy'], label='val_accuracy')
plt.legend(loc='lower right')
plt.title('accuracy wo dropout')
plt.xlabel('epoch')
plt.ylabel('accuracy')
```

```
plt.subplot(232)
plt.plot(history_dropout_hidden.history['accuracy'], label='accuracy')
plt.plot(history_dropout_hidden.history['val_accuracy'], label='val_accuracy')
plt.legend(loc='lower right')
plt.title('accuracy w dropout connected to input')
plt.xlabel('epoch')
plt.ylabel('accuracy')
```

```
plt.subplot(233)
plt.plot(history_dropout_hidden2.history['accuracy'], label='accuracy')
plt.plot(history_dropout_hidden2.history['val_accuracy'], label='val_accuracy')
plt.legend(loc='lower right')
plt.title('accuracy w dropouts between hidden layers')
plt.xlabel('epoch')
plt.ylabel('accuracy')
#### Fill in plot ####
```

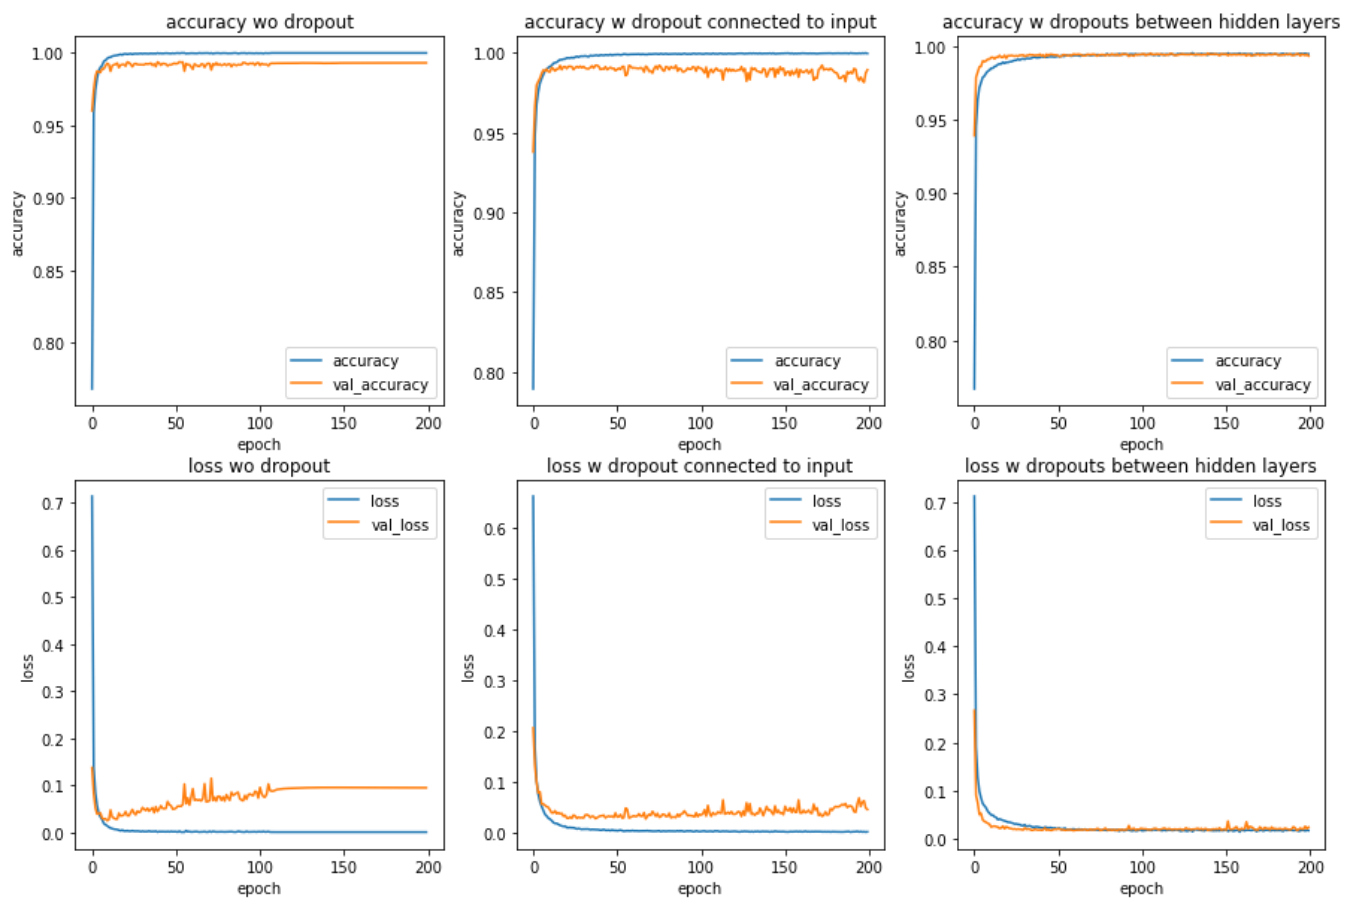
```
#Plot loss vs epoch
```

```
plt.subplot(234)
plt.plot(history_overfit.history['loss'], label='loss')
plt.plot(history_overfit.history['val_loss'], label='val_loss')
plt.legend(loc='upper right')
plt.title('loss wo dropout')
plt.xlabel('epoch')
plt.ylabel('loss')
```

```
plt.subplot(235)
plt.plot(history_dropout_hidden.history['loss'], label='loss')
plt.plot(history_dropout_hidden.history['val_loss'], label='val_loss')
plt.legend(loc='upper right')
plt.title('loss w dropout connected to input')
plt.xlabel('epoch')
plt.ylabel('loss')
```

```
plt.subplot(236)
plt.plot(history_dropout_hidden2.history['loss'], label='loss')
plt.plot(history_dropout_hidden2.history['val_loss'], label='val_loss')
plt.legend(loc='upper right')
plt.title('loss w dropouts between hidden layers')
plt.xlabel('epoch')
plt.ylabel('loss')
#### Fill in plot ####
```

```
plt.show()
```



## Section 2- Autoencoders

### 2.1 Linear AE

Fill in the model:

- Input: Flattened grayscale image to  $28^2 = 784$ -dimensional vector.
- 1st hidden: 400 perceptrons.
- 2nd hidden: 200 perceptrons.
- 3rd hidden: 100 perceptrons.
- 4th hidden: 200 perceptrons.
- 5th hidden: 400 perceptrons.
- Output: 784 perceptrons.

**Train for 150 epochs**

```
In [0]: #Reshape training and testing data
data_train_reshape_fcae = data_train.reshape(data_train.shape[0], 784)
data_test_reshape_fcae = data_test.reshape(data_test.shape[0], 784)
```

```

In [154]: # Create autoencoder architecture
def deep_ae():
    model = tf.keras.models.Sequential()

    model.add(tf.keras.layers.Input(shape=(784,)))
    # Encoder

    ##### Fill in the model #####
    model.add(tf.keras.layers.Dense(400, activation='relu'))
    #model.add(tf.keras.layers.Dropout(0.3))
    model.add(tf.keras.layers.Dense(200, activation='relu'))
    #model.add(tf.keras.layers.Dropout(0.3))
    model.add(tf.keras.layers.Dense(100, activation='relu'))
    model.add(tf.keras.layers.Dropout(0.3))
    # Decoder

    ##### Fill in the model #####
    model.add(tf.keras.layers.Dense(200, activation='relu'))
    model.add(tf.keras.layers.Dropout(0.3))
    model.add(tf.keras.layers.Dense(400, activation='relu'))
    model.add(tf.keras.layers.Dropout(0.3))
    model.add(tf.keras.layers.Dense(784, activation='sigmoid'))

    return model

#Create deep autoencoder graph, compile it to use mean squared error loss and the adam optimizer,
train the model, create predictions
deep_ae = deep_ae()
print(deep_ae.summary())
deep_ae.compile(loss='categorical_crossentropy', optimizer='adam')
history_deep_ae = deep_ae.fit(data_train_reshape_fcae, data_train_reshape_fcae, validation_data=(data_test_reshape_fcae, data_test_reshape_fcae), epochs=10, batch_size=250, shuffle=True)
decoded_data = deep_ae.predict(data_test_reshape_fcae)

#Obtain encoder representation of data
get_hl = K.function([deep_ae.layers[0].input], [deep_ae.layers[2].output])
deep_ae_hl = get_hl([data_test_reshape_fcae])[0]

```

Model: "sequential\_30"

Layer (type)	Output Shape	Param #
dense_24 (Dense)	(None, 400)	314000
dense_25 (Dense)	(None, 200)	80200
dense_26 (Dense)	(None, 100)	20100
dropout_21 (Dropout)	(None, 100)	0
dense_27 (Dense)	(None, 200)	20200
dropout_22 (Dropout)	(None, 200)	0
dense_28 (Dense)	(None, 400)	80400
dropout_23 (Dropout)	(None, 400)	0
dense_29 (Dense)	(None, 784)	314384

Total params: 829,284  
Trainable params: 829,284  
Non-trainable params: 0

None

Epoch 1/10

240/240 [=====] - 1s 4ms/step - loss: 575.7485 - val\_loss: 554.7593

Epoch 2/10

240/240 [=====] - 1s 4ms/step - loss: 548.4683 - val\_loss: 544.8527

Epoch 3/10

240/240 [=====] - 1s 4ms/step - loss: 542.5052 - val\_loss: 539.8880

Epoch 4/10

240/240 [=====] - 1s 4ms/step - loss: 539.5171 - val\_loss: 537.9699

Epoch 5/10

240/240 [=====] - 1s 4ms/step - loss: 537.7516 - val\_loss: 535.8461

Epoch 6/10

240/240 [=====] - 1s 4ms/step - loss: 536.5424 - val\_loss: 534.6545

Epoch 7/10

240/240 [=====] - 1s 4ms/step - loss: 535.5967 - val\_loss: 533.7915

Epoch 8/10

240/240 [=====] - 1s 4ms/step - loss: 534.9063 - val\_loss: 533.5900


Epoch 9/10

240/240 [=====] - 1s 4ms/step - loss: 534.4302 - val\_loss: 532.9552

Epoch 10/10

240/240 [=====] - 1s 4ms/step - loss: 533.9613 - val\_loss: 532.4114


```
In [0]: #Plot train/validation loss vs epoch
plt.figure(figsize=(15,10))
#Plot loss vs epoch
plt.plot(history_deep_ae.history['loss'], label='loss')
plt.plot(history_deep_ae.history['val_loss'], label='val_loss')
plt.legend(loc='upper right')
plt.title('loss')
plt.xlabel('epoch')
plt.ylabel('loss')
#### Fill in the plot ####
plt.show()
```



```
In [0]: #Plot samples of 10 images, their hidden layer representations, and their reconstructions
n = 10 # how many digits we will display
plt.figure(figsize=(20, 4))
for i in range(n):
    # display original
    ax = plt.subplot(3, n, i + 1)
    plt.imshow(data_test_reshape_fcae[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # display hidden layer representation
    ax = plt.subplot(3, n, i + 1 + n)
    plt.imshow(deep_ae_hl[i].reshape(10, 10))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # display reconstruction
    ax = plt.subplot(3, n, i + 1 + n + n)
    plt.imshow(decoded_data[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```



## 2.2 Convolutional AE

Fill in the model:

- Input: 28x28x1 grayscale image.
- 1st hidden: 2D convolutional layer with 16 feature maps and 3x3 filters.
- 2nd hidden: A 2x2 maxpool layer.
- 3rd hidden: 2D convolutional layer with 8 feature maps and 3x3 filters.
- 4th hidden: A 2x2 maxpool layer.
- 5th hidden: 2D convolutional layer with 8 feature maps and 3x3 filters.
- 6th hidden: A 2x2 upsample layer.
- 7th hidden: 2D convolutional layer with 16 feature maps and 3x3 filters.
- 8th hidden: A 2x2 upsample layer.
- Output: A convolutional layer with a single feature map and 3x3 filters.

**All experiments with dropout set at 30%. Train for 200 epochs**



```
In [0]: #Reshape data to account for grayscale channel in each image  
data_train_reshape_cae = data_train.reshape(data_train.shape[0], 28, 28, 1)  
data_test_reshape_cae = data_test.reshape(data_test.shape[0], 28, 28, 1)
```

```

In [157]: #Create Convolutional AutoEncoder Architecture
def cae():
    model = tf.keras.models.Sequential()

    #Encoder
    model.add(tf.keras.layers.Conv2D(16, kernel_size=(3, 3), activation='relu', padding='same', input_shape=(28, 28, 1)))
    model.add(tf.keras.layers.MaxPooling2D(pool_size=(2, 2), padding='same'))
    model.add(tf.keras.layers.Conv2D(8, (3, 3), activation='relu', padding='same'))
    model.add(tf.keras.layers.MaxPooling2D(pool_size=(2, 2), padding='same'))
    #### Fill in model ####

    #Decoder
    model.add(tf.keras.layers.Conv2D(8, (3, 3), activation='relu', padding='same'))
    model.add(tf.keras.layers.UpSampling2D((2, 2)))
    model.add(tf.keras.layers.Conv2D(16, (3, 3), activation='relu', padding='same'))
    model.add(tf.keras.layers.UpSampling2D((2, 2)))
    #### Fill in model ####
    model.add(tf.keras.layers.Conv2D(1, (3, 3), activation='relu', padding='same'))

    return model

conv_ae = cae()
print(conv_ae.summary())

#Create deep autoencoder graph, compile it to use mean squared error loss and the adam optimizer, train the model, create predictions
conv_ae.compile(loss='mse', optimizer='adam')
history_conv_ae = conv_ae.fit(data_train_reshape_cae, data_train_reshape_cae, validation_data=(data_test_reshape_cae, data_test_reshape_cae), epochs=10, batch_size=250, shuffle=True)
decoded_data = conv_ae.predict(data_test_reshape_cae)

#Obtain encoder representation of data
get_hl = K.function([conv_ae.layers[0].input], [conv_ae.layers[3].output])
conv_ae_hl = get_hl([data_test_reshape_cae])[0]

```

Model: "sequential\_32"

Layer (type)	Output Shape	Param #
conv2d_151 (Conv2D)	(None, 28, 28, 16)	160
max_pooling2d_68 (MaxPooling)	(None, 14, 14, 16)	0
conv2d_152 (Conv2D)	(None, 14, 14, 8)	1160
max_pooling2d_69 (MaxPooling)	(None, 7, 7, 8)	0
conv2d_153 (Conv2D)	(None, 7, 7, 8)	584
up_sampling2d_58 (UpSampling)	(None, 14, 14, 8)	0
conv2d_154 (Conv2D)	(None, 14, 14, 16)	1168
up_sampling2d_59 (UpSampling)	(None, 28, 28, 16)	0
conv2d_155 (Conv2D)	(None, 28, 28, 1)	145

Total params: 3,217

Trainable params: 3,217

Non-trainable params: 0

None

Epoch 1/10

240/240 [=====] - 2s 9ms/step - loss: 0.0254 - val\_loss: 0.0102

Epoch 2/10

240/240 [=====] - 2s 8ms/step - loss: 0.0089 - val\_loss: 0.0077

Epoch 3/10

240/240 [=====] - 2s 8ms/step - loss: 0.0074 - val\_loss: 0.0067

Epoch 4/10

240/240 [=====] - 2s 8ms/step - loss: 0.0066 - val\_loss: 0.0061

Epoch 5/10

240/240 [=====] - 2s 8ms/step - loss: 0.0060 - val\_loss: 0.0056

Epoch 6/10

240/240 [=====] - 2s 8ms/step - loss: 0.0056 - val\_loss: 0.0054

Epoch 7/10

240/240 [=====] - 2s 8ms/step - loss: 0.0053 - val\_loss: 0.0051

Epoch 8/10

240/240 [=====] - 2s 8ms/step - loss: 0.0051 - val\_loss: 0.0049

Epoch 9/10


240/240 [=====] - 2s 9ms/step - loss: 0.0049 - val\_loss: 0.0047

Epoch 10/10

240/240 [=====] - 2s 8ms/step - loss: 0.0048 - val\_loss: 0.0046

```
In [0]: #Plot train/validation loss vs epoch  
#Plot loss vs epoch  
plt.figure(figsize=(15,10))  
plt.plot(history_conv_ae.history['loss'], label='loss')  
plt.plot(history_conv_ae.history['val_loss'], label='val_loss')  
plt.legend(loc='upper right')  
plt.title('loss')  
plt.xlabel('epoch')  
plt.ylabel('loss')  
plt.plot()  
#### Fill in the plot ####
```

Out[0]: []




```
In [0]: #Plot samples of 10 images, their hidden layer representations, and their reconstructions
n = 10 # how many digits we will display
plt.figure(3)
plt.figure(figsize=(20, 4))
for i in range(n):
    # display original
    ax = plt.subplot(3, n, i + 1)
    plt.imshow(data_test_reshape_cae[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # display hidden layer representation
    ax = plt.subplot(3, n, i + 1 + n)
    plt.imshow(conv_ae_hl[i].reshape(28, 14))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # display reconstruction
    ax = plt.subplot(3, n, i + 1 + n + n)
    plt.imshow(decoded_data[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```

<Figure size 432x288 with 0 Axes>



## 2.3 Machine Anomaly Detection

At this point you have enough starter code to Using the dataset provided **create the autoencdoer model** you deem necessary to achieve better than 75 true positives (TP = 75) where a true instance is an anomaly. Or detect all 143 if you can! Although anomaly detection thresholds can be set arbitrarily and various metrics are used depending on the problem, we will set ours at 2 standard deviations from the mean of "normal" data to judge TP's. Use the code provided at the bottom for calculating true positives and histogramming.

```
In [0]: ##### Restart your kernal and run from here to clear some memory
import numpy as np
import matplotlib.pyplot as plt

import tensorflow as tf
from tensorflow import keras
tf.keras.backend.set_floatx('float64')

import sys
from os import listdir
from os.path import isfile, join
```

```
In [0]: ### I am using colab and the data is saved in my google drive so I need:
print('attention! using colab here and data is in my google drive!')
from google.colab import drive
drive.mount('/content/drive')
```

attention! using colab here and data is in my google drive!

Go to this URL in a browser: [https://accounts.google.com/o/oauth2/auth?client\\_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect\\_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aob&response\\_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly](https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aob&response_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly)

Enter your authorization code:

• • • • •


Mounted at /content/drive

```
In [0]: datapath = './drive/My Drive/UCSD/2020/ECE228/PY2/'
```

## Example spectrograms

```
In [0]: ##### Load melspectrograms
ex_norm = np.load(datapath + 'ex_normalspec.npy')
ex_anom = np.load(datapath + 'ex_abnormspec.npy')
plt.figure(figsize=(15,7))
plt.subplot(211)
plt.imshow(ex_norm[0,::-1], origin='lower', cmap='inferno')
plt.xlabel('Time')
plt.ylabel('Mel Filter Number')
plt.title('Normal Data Melspectrogram (64 Mels)')

plt.subplot(212)
plt.imshow(ex_anom[0,::-1], origin='lower', cmap='inferno')
plt.xlabel('Time')
plt.ylabel('Mel Filter Number')
plt.title('Anomalous Data Melspectrogram (64 Mels)')
plt.show()
```



### General template, up to this point, for constructing your deep learning model

1. Set up the data (reshape, scale, etc...
2. Initialize a loss function
3. Compile a model
4. Train a model

```

In [142]: # Create your own Baseline autoencoder
# Model name is fixed for use by later code
autoencoderBASE = tf.keras.models.Sequential([
    #### Fill in your model ####
    keras.layers.Input(shape=(64, 312, 1)),

    keras.layers.Conv2D(32, kernel_size=(3, 3), activation='relu', padding='same', kernel_initializer=
'random_normal'),
    keras.layers.MaxPooling2D(pool_size=(2, 2), padding='same'),
    keras.layers.Dropout(0.3),
    keras.layers.Conv2D(16, kernel_size=(3, 3), activation='relu', padding='same', kernel_initializer=
'random_normal'),
    keras.layers.MaxPooling2D(pool_size=(2, 2), padding='same'),
    keras.layers.Conv2D(4, (3, 3), activation='relu', padding='same', kernel_initializer='random_norm
al'),
    keras.layers.MaxPooling2D(pool_size=(2, 2), padding='same'),

    keras.layers.Conv2D(4, (3, 3), activation='relu', padding='same', kernel_initializer='random_norm
al'),
    keras.layers.UpSampling2D((2, 2)),
    keras.layers.Conv2D(16, kernel_size=(3, 3), activation='relu', padding='same', kernel_initializer=
'random_normal'),
    keras.layers.UpSampling2D((2, 2)),
    keras.layers.Conv2D(32, (3, 3), activation='relu', padding='same', kernel_initializer='random_nor
mal'),
    keras.layers.UpSampling2D((2, 2)),
    #keras.layers.Dropout(0.3),

    keras.layers.Conv2D(1, (3, 3), activation='relu', padding='same', kernel_initializer='random_norm
al'),
])
print(autoencoderBASE.summary())

```



Model: "sequential\_23"

Layer (type)	Output Shape	Param #
conv2d_131 (Conv2D)	(None, 64, 312, 32)	320
max_pooling2d_55 (MaxPooling)	(None, 32, 156, 32)	0
dropout_6 (Dropout)	(None, 32, 156, 32)	0
conv2d_132 (Conv2D)	(None, 32, 156, 16)	4624
max_pooling2d_56 (MaxPooling)	(None, 16, 78, 16)	0
conv2d_133 (Conv2D)	(None, 16, 78, 4)	580
max_pooling2d_57 (MaxPooling)	(None, 8, 39, 4)	0
conv2d_134 (Conv2D)	(None, 8, 39, 4)	148
up_sampling2d_53 (UpSampling)	(None, 16, 78, 4)	0
conv2d_135 (Conv2D)	(None, 16, 78, 16)	592
up_sampling2d_54 (UpSampling)	(None, 32, 156, 16)	0
conv2d_136 (Conv2D)	(None, 32, 156, 32)	4640
up_sampling2d_55 (UpSampling)	(None, 64, 312, 32)	0
conv2d_137 (Conv2D)	(None, 64, 312, 1)	289

Total params: 11,193  
Trainable params: 11,193  
Non-trainable params: 0

None

```
In [0]: # Load data
x_train = np.load(datapath + 'training_data.npy')
anomaly_data = np.load(datapath + 'test_data.npy')

x_train = x_train[:,0,:,-1]
anomaly_data = anomaly_data[:,0,:,-1]
```

```
In [0]: # Normal
#x_train = ( x_train - np.amin(x_train) )/( np.amax(x_train) - np.amin(x_train) )
#anomaly_data = ( anomaly_data - np.amin(anomaly_data) )/( np.amax(anomaly_data) - np.amin(anomaly_data) )
test = anomaly_data


# reshape
normaldata_reshape = x_train.reshape(x_train.shape[0],64,312,1)
test_reshape = test.reshape(test.shape[0],64,312,1)

x_train_reshape = normaldata_reshape[:800,:,:,:]
x_val_reshape = normaldata_reshape[800,:,:,:]
```

```
In [158]: # Training
autoencoderBASE.compile(loss='mse', optimizer='adam')
history_conv_BASE = autoencoderBASE.fit(x_train_reshape, x_train_reshape, validation_data=(x_val_reshape, x_val_reshape), epochs=10, batch_size=50, shuffle=True)
decoded_data = autoencoderBASE.predict(test_reshape)
```

```
Epoch 1/10
16/16 [=====] - 1s 64ms/step - loss: 12.1430 - val_loss: 11.2716
Epoch 2/10
16/16 [=====] - 1s 57ms/step - loss: 10.0792 - val_loss: 10.4480
Epoch 3/10
16/16 [=====] - 1s 58ms/step - loss: 9.7582 - val_loss: 10.8751
Epoch 4/10
16/16 [=====] - 1s 58ms/step - loss: 9.6825 - val_loss: 10.7301
Epoch 5/10
16/16 [=====] - 1s 58ms/step - loss: 9.6608 - val_loss: 10.6104
Epoch 6/10
16/16 [=====] - 1s 58ms/step - loss: 9.6490 - val_loss: 10.5816
Epoch 7/10
16/16 [=====] - 1s 57ms/step - loss: 9.6391 - val_loss: 10.5560
Epoch 8/10
16/16 [=====] - 1s 58ms/step - loss: 9.6334 - val_loss: 10.5583
Epoch 9/10
16/16 [=====] - 1s 58ms/step - loss: 9.6299 - val_loss: 10.5252
Epoch 10/10
16/16 [=====] - 1s 58ms/step - loss: 9.6237 - val_loss: 10.5395
```

```
In [144]: # Plot loss versus epoch.
plt.figure(figsize=(15, 7))
plt.plot(history_conv_BASE.history['loss'], label='loss')
plt.plot(history_conv_BASE.history['val_loss'], label='val_loss')
plt.legend(loc='upper right')
plt.title('loss')
plt.xlabel('epoch')
plt.ylabel('loss')
plt.show()
```



```

In [138]: ##### This code should remain untouched as much as possible,
##### except where your variable names for loss function or data set are needed.
##### This code feeds your data through the trained network to get mean and std
##### If you did not use a validation set then only use
##### your training data. Concatenating is therefore un-needed.
##### lossFunction <- Your loss function's name or use this one. Your choice of loss function.
#####

lossFunction = tf.keras.losses.MeanSquaredError()

norm_list = []
dataset = (tf.data.Dataset.from_tensor_slices(normaldata_reshape)).batch(1)
for i, instance in dataset.enumerate():
    ae_predictions = autoencoderBASE(instance).numpy()
    norm_list.append(lossFunction(instance, ae_predictions).numpy())
# Feed the anomaly data through to get its error
anom_list = []
anomset = (tf.data.Dataset.from_tensor_slices(test_reshape)).batch(1)
for i, instance in anomset.enumerate():
    ae_predictions = autoencoderBASE(instance).numpy()
    anom_list.append(lossFunction(instance, ae_predictions).numpy())

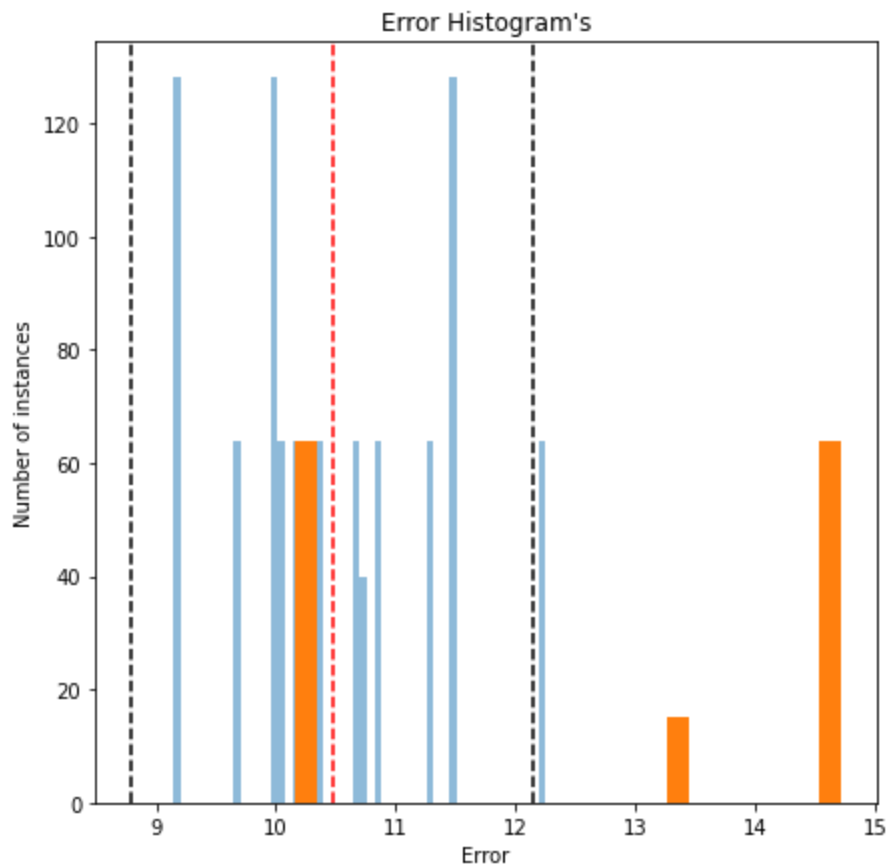
normal_data_ERRORS = np.array(norm_list)
abnormal_data_ERRORS = np.array(anom_list)

##### Code for presenting true positives to Question 2.3 #####
threshold = 2
mean = normal_data_ERRORS.mean()
std = normal_data_ERRORS.std()
print(f'The mean of normal data is {mean:.4f} \
      and standard deviation is {std:.4f}')
upperbound = mean+threshold*std
lowerbound = mean-threshold*std
plt.figure(figsize=(7,7))
plt.title('Error Histogram\ 's')
plt.hist(normal_data_ERRORS, bins=50, alpha=0.5)
plt.hist(abnormal_data_ERRORS, bins=25, alpha=1.0)
plt.axvline(mean, ls='--', c='r')
plt.axvline(lowerbound, ls='--', c='k')
plt.axvline(upperbound, ls='--', c='k')
plt.xlabel('Error')
plt.ylabel('Number of instances')
plt.show()

```

The mean of normal data is 10.4711

and standard deviation is 0.8402



```
In [139]: tp_count = np.sum(abnormal_data_ERRORS >= upperbound) +\
           np.sum(abnormal_data_ERRORS <= lowerbound)
fn_count = anomaly_data.shape[0] - tp_count
fp_count = np.sum(normal_data_ERRORS >= upperbound) +\
           np.sum(normal_data_ERRORS <= lowerbound)
tn_count = 1000 - fp_count

print(f' TP {tp_count} \t FP {fp_count}')
```

```
print(f' FN {fn_count} \t TN {tn_count}')
```

TP 79	FP 64
FN 64	TN 936